

OK I see your current Vercel and Render values. Both are **missing a lot** — especially the LLM API keys for Jarvis, the tiered model setup, Paddle, CORS/CSRF, and more. Let me also check what the frontend actually needs for the Jarvis/onboarding flow:
Now I can see exactly what the frontend needs. Let me also check the Smart Router to confirm the model names:
Now I have everything I need. Let me create the complete env var lists for both platforms:
Here are the **COMPLETE** values you need to paste. Copy-paste each one exactly.

VERCEL (Frontend) — All Environment Variables

...

```
NEXT_PUBLIC_API_URL=https://parwa-backend.onrender.com
NEXT_PUBLIC_WS_URL=wss://parwa-backend.onrender.com
NEXT_PUBLIC_GOOGLE_CLIENT_ID=810864777209-tpgn9avs8evtin6kjavgr3s162kp8ju0a.app
s.googleusercontent.com
NEXT_PUBLIC_PADDLE_KEY=live_84ceb40f4a03f934aadd1460d60
NEXT_PUBLIC_JWT_ALGORITHM=HS256
NEXT_PUBLIC_APP_VERSION=0.1.0
NEXT_PUBLIC_ENV=production
NEXTAUTH_SECRET=sDdNJM9/HDJZzn9XDhlaob9KK6qJ5LZoMRNLxR4ohN7NZkDPvj4KIO
Vc+Q+mC34N
NEXTAUTH_URL=https://parwa.buzz
GOOGLE_CLIENT_ID=810864777209-tpgn9avs8evtin6kjavgr3s162kp8ju0a.apps.googleusercon
tent.com
GOOGLE_CLIENT_SECRET=GOCS PX-XINil3P8qxGN4HCyacZol_YEzMI
SERVER_API_URL=https://parwa-backend.onrender.com
JWT_SECRET_KEY=parwa_jwt_sk_2026_aZ7bX3cV5dN8mK2pQ9wR4tY6ul0oL
GOOGLE_AI_API_KEY=AlzaSyATHbcolmlaNufj6ZHR6tebMmlqqcmCsEs
GROQ_API_KEY=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
CEREBRAS_API_KEY=csk-dvej6d3jvyw5nc4r53hy2k5kp543d9vx5yt8jtr2yfffwwwvh
```

...

Why the LLM keys on Vercel?

Your **Jarvis chat** (`/api/jarvis/[...path]/route.ts`) and **chat completions** (`/api/chat/completions/route.ts`) and **onboarding Jarvis** all use these API keys directly as a **fallback** when the backend is down or sleeping. They call Google AI / Groq / Cerebras from the Vercel serverless function itself. Without these keys, Jarvis won't work when the backend is cold-starting.

RENDER (Backend) — All Environment Variables

...

ENVIRONMENT=production
SECRET_KEY=parwa_dev_sk_2026_xK9mP2vR8wQ4tY6ul0oL3jH7nB5cF1dE
DEBUG=false
APP_VERSION=0.1.0
DATABASE_URL=postgresql://postgres.fmpibdauppnzsfisodkhp:Durgamaa@754@aws-1-ap-northeast-1.pooler.supabase.com:6543/postgres
JWT_SECRET_KEY=parwa_jwt_sk_2026_aZ7bX3cV5dN8mK2pQ9wR4tY6ul0oL
JWT_ALGORITHM=HS256
JWT_ACCESS_TOKEN_EXPIRE_MINUTES=15
JWT_REFRESH_TOKEN_EXPIRE_DAYS=7
LLM_PROVIDER=litellm
LLM_PRIMARY_PROVIDER=google
LLM_FALLBACK_PROVIDER=groq
GOOGLE_AI_API_KEY=AlzaSyATHbcolmlaNufj6ZHR6tebMmlqqcmCsEs
CEREBRAS_API_KEY=csk-dvej6d3jvyw5nc4r53hy2k5kp543d9vx5yt8jtr2yfffwwvh
GROQ_API_KEY=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
AI_LIGHT_MODEL=gemini/gemini-2.0-flash-lite
AI_MEDIUM_MODEL=gemini/gemini-2.0-flash
AI_HEAVY_MODEL=groq/llama-3.3-70b-versatile
AI_FAILOVER_MODEL=cerebras/llama-4-scout-17b-16e-instruct
GOOGLE_CLIENT_ID=810864777209-tpgn9avs8evtin6kjavgr3s162kp8ju0a.apps.googleusercontent.com
GOOGLE_CLIENT_SECRET=GOCSPX-XINil3P8qxGN4HCyacZol_YEzMI
PADDLE_CLIENT_TOKEN=live_84ceb40f4a03f934aadd1460d60
PADDLE_API_KEY=pdl_live_apikey_01kkze0w44439v7e8s864xgpm5_VdF7WkCyyR3JSBGw62t2VN_A20
PADDLE_WEBHOOK_SECRET=pdl_ntfset_01kqphdj5g7338706wyqbxbaq3_hAwfyiM2+o6S62a16y8cmfTjk7Tc5JnL
PADDLE_WEBHOOK_NOTIFICATION_SET_ID=ntfset_01kqphdj5g7338706wyqbxbaq3
PADDLE_PRICE_IDS={"mini_parwa":"pri_01krxm4r0kcm6mm5fc84pp9bj0","parwa":"pri_01krxm4ra529ry7bzt9z73pza1","parwa_high":"pri_01krxm4rjx1bfgg1w9z4qr3dd8"}
NEXT_PUBLIC_PADDLE_KEY=live_84ceb40f4a03f934aadd1460d60
DATA_ENCRYPTION_KEY=parwa_enc_key_2026_xK9mP2vR8w0123456789012
PRICING_SIGNING_KEY=parwa_pricing_sk_2026_mK2pQ9wR4tY6ul0oL3jH7n
REFRESH_TOKEN_PEPPEER=parwa_refresh_pepper_dev_2026
CORS_ORIGINS=https://parwafrontend.vercel.app,https://parwa.buzz
CSRF_TRUSTED_ORIGINS=https://parwafrontend.vercel.app,https://parwa.buzz
FRONTEND_URL=https://parwa.buzz
NEXTAUTH_SECRET=uH3196lsyUVQ_Gvldx5W-dR3c0GKzAWL3eh1uW7GbBlyxzuofrykb85nzNjfqzHb
FROM_EMAIL=noreply@parwa.ai
GDPR_RETENTION_DAYS=365

```
AUDIT_LOG_RETENTION_DAYS=2555
IP_ALLOWLIST_ENABLED=false
'''
```

How the LLM Tiers Work (Your 3 Keys → 11 Models)

Your **3 API keys** power **11 models** across **4 tiers** with automatic failover:

Tier	Models (Priority Order)	Used For	API Key Used
LIGHT (90% of calls)	Cerebras Llama 3.1 8B → Groq Llama 3.1 8B → Google Gemma 3 27B	Intent, PII, sentiment, classification, GSD states, simple drafts	Cerebras → Groq → Google
MEDIUM (8%)	Google Gemini 3.1 Flash-Lite → Google Gemini 2.5 Flash → Groq Llama 3.3 70B → Groq Qwen3 32B	Moderate/complex responses, reasoning, reflexion	Google → Groq → Groq
HEAVY (2%)	Groq GPT-OSS 120B → Cerebras GPT-OSS 120B → Groq Llama 4 Scout	Complex drafts, deep reflexion cycles	Groq → Cerebras → Groq
GUARDRAIL	Groq Llama Guard 4 12B	Safety checks on all responses	Groq

LiteLLM auto-routes `cerebras/`, `groq/`, `gemini/` prefixes to the correct API key. The Smart Router handles priority-based failover — if the primary model in a tier fails, it automatically tries the next one.

- ### Variant Access (what each plan gets):
- **Mini PARWA** → LIGHT + GUARDRAIL only
 - **PARWA** → LIGHT + MEDIUM + GUARDRAIL
 - **PARWA High** → All tiers (LIGHT + MEDIUM + HEAVY + GUARDRAIL)